

SHUSHANT KAMATAR

✉ shushantkamatar195@gmail.com ☎ +91-8088348874 LinkedIn Github 🌐 Coding Profile

Experience

Forin Agency

Machine Learning Engineer

Nov 2025 – Present

- Architected a high-concurrency real-time **Interview Intelligence System** using **FastAPI** and **WebSockets** to enable full-duplex communication, achieving **< 800ms latency** by parallelizing audio transcription (**Deepgram**) and LLM inference to support seamless, lag-free live interactions.
- Engineered a context-aware **Retrieval-Augmented Generation (RAG)** with **Supabase (pgvector)** power a live recommendation engine to dynamically retrieve candidate resume data and historical "winning questions" during live calls, enhancing interviewer decision-making speed.
- Developed a secure, production-grade Chrome Extension utilizing the specialized '**Offscreen**' API to capture high-fidelity system audio streams while effectively bypassing standard browser sandbox restrictions, ensuring reliable audio data ingestion for the inference engine.

Lenovo

Software Engineer

May 2024 – Sep 2025

- Engineered a robust real-time **NLP Inference Pipeline** using **Python** and **Flask** to continuously stream and process sentiment analysis on high-volume product feedback, effectively transforming raw unstructured text into live, actionable strategic intelligence for product stakeholders.
- Optimized critical legacy ETL pipelines by re-architecting iterative Python loops into efficient **vectorized Pandas operations**, reducing processing time by **99%** (slashing execution windows from multiple hours to mere seconds) and enabling **real-time analytics** for the leadership team.
- Orchestrated the complex technical deployment of data-driven features across **Western Europe (15 markets)**, meticulously managing diverse localization requirements and ensuring **system reliability** during high-traffic regional rollouts and peak user activity periods.

Skills

- Languages:** Python, JavaScript, SQL, HTML/CSS
- Frameworks:** FastAPI, Flask, PyTorch, Deepgram, LangChain
- Tools:** Docker, Kubernetes, Github Actions, MLflow, DVC, Prometheus
- Machine Learning & AI:** Sciketlearn, NLP, RAG, Vector Embeddings
- System Design & Backend:** AsyncIO, Redis, WebSockets, Supabase, FAISS

Education

RV College of Engineering

B.E in Aerospace Engineering; CGPA: 7.07

Jun 2019 – Aug 2023

Bangalore Urban

Projects

Enterprise Natural Language-to-SQL Engine

Dec 2025

- Architected a semantic schema retrieval system using **FAISS** to dynamically inject relevant database context into the LLM context window, reducing token usage by 90% and enabling robust support for complex schemas (**100+ tables**) by eliminating hallucination risks.
- Migrated inference stack from GPT-4 to **Llama 3.3** on **Groq LPUs**, reducing query latency to **<200ms** and operational costs by **75%** while maintaining SQL syntax accuracy through the implementation of advanced prompt distillation techniques and rigorous regression testing.
- Engineered a deterministic **Self-Correction Loop** that autonomously parses database execution errors (e.g., 'no such column' or syntax faults) to iteratively repair and re-execute malformed SQL queries, successfully recovering **40%** of failures without human intervention.

Tech Stack: SQLite, Python, Redis, Docker, FAISS, Streamlit

High-Throughput Inference Gateway

Nov 2025

- Developed a resilient, unified API gateway using **FastAPI** and **Async Queues** to efficiently route concurrent requests across multiple heterogenous LLM providers, ensuring high availability through the implementation of automated failover strategies and intelligent load balancing .
- Implemented a hybrid semantic caching layer combining **Redis** and **Vector Embeddings** to intercept and serve recurring natural language queries , improving **P99** latency for recurring queries by **600ms** and reducing external API costs by **30%** by serving responses directly from the cache.

Tech Stack: FastAPI, Python, Redis, Docker

Certificate

- Nasscom Certified Data Scientist 